




Image Classification using **K-Nearest Neighbor Learning**

MNIST Database



NCSSM CS4320 Machine Learning
Srihas Surapaneni




01

Image Classification

Problem Description



This technology, **optical character recognition (OCR)**, is incredibly useful for data processing, accessibility, and automation.



02

MNIST


Modified National Institute of Standards and Technology

MNIST CSV Dataset

MNIST is a widely used dataset of 70,000 handwritten digits from 0 to 9, found on [Kaggle](#).

MNIST CSV reformats each **image** as a **row of data**.

- 60,000 Train + 10,000 Test
- Each Row is 785 Columns
 - Label + 784 Values
 - 28×28 Grid of Values
- Each Value is Between 0-255
 - 0 - White
 - 255 - Black




MNIST CSV Dataset

MNIST is a widely used dataset of 70,000 handwritten digits from 0 to 9, found on [Kaggle](#).

MNIST CSV reformats each image as a row of data.

- 60,000 Train + 10,000 Test
- Each Row is 785 Columns
 - Label + 784 Values
 - 28x28 Grid of Values
- Each Value is Between 0-255
 - 0 - White
 - 255 - Black

MNIST Digit: 5




MNIST CSV Dataset

MNIST is a widely used dataset of 70,000 handwritten digits from 0 to 9, found on [Kaggle](#).

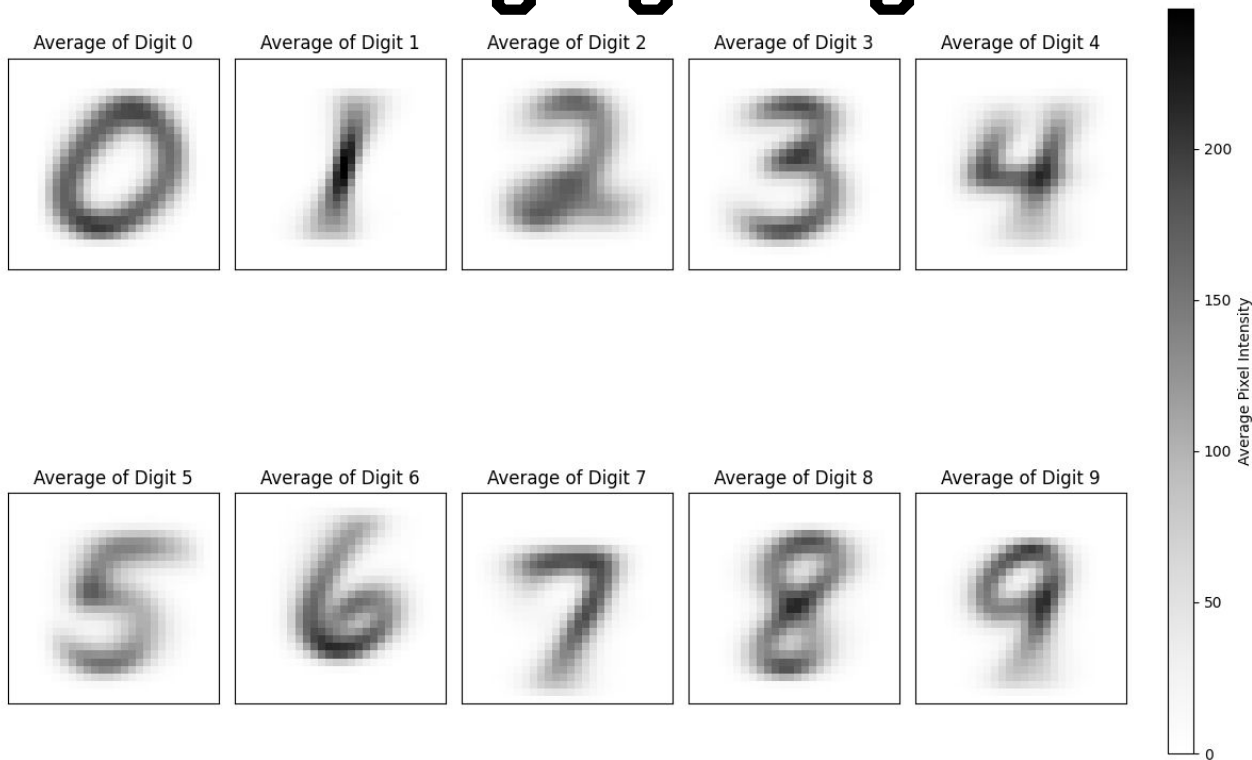
MNIST CSV reformats each image as a row of data.

- 60,000 Train + 10,000 Test
- Each Row is 785 Columns
 - Label + 784 Values
 - 28x28 Grid of Values
- Each Value is Between 0-255
 - 0 - White
 - 255 - Black






Averaging Images



We can visualize an approximate of what KNN is testing against by averaging the **intensity** of every **pixel** for every **digit**, and produce the **"average image"** of each digit.



03

K-Nearest Neighbors

Distance Based Classification Algorithm



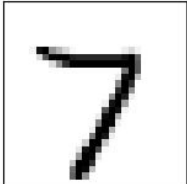
Why KNN


What is KNN?

- Uses reference values to calculate numerical distances between test and training data
- Uses the “k” points in the training data to classify a test point.

Why KNN is perfect for this problem:

- Data is already standardized from 0-255.
- Non-binary classification so referencing training data increases accuracy.

Test Point True: 7 Predicted: 7	Neighbor 1 Label: 7 Dist: 676.58	Neighbor 2 Label: 7 Dist: 793.99	Neighbor 3 Label: 7 Dist: 862.68	Neighbor 4 Label: 7 Dist: 864.50	Neighbor 5 Label: 7 Dist: 894.70
					



04


Model Validation

Training, Validation, and Testing Methods




Misclassification


Test Example (idx: 33)
True: 4
Predicted: 0 (Misclassified)




Neighbor 1
Label: 4
Dist: 1384.31




Neighbor 2
Label: 0
Dist: 1780.64




Neighbor 3
Label: 4
Dist: 1841.15




Neighbor 4
Label: 0
Dist: 1897.24




Neighbor 5
Label: 0
Dist: 1904.74




Test Example (idx: 115)
True: 4
Predicted: 9 (Misclassified)




Neighbor 1
Label: 9
Dist: 1126.15




Neighbor 2
Label: 9
Dist: 1140.99




Neighbor 3
Label: 9
Dist: 1142.65




Neighbor 4
Label: 9
Dist: 1150.51




Neighbor 5
Label: 9
Dist: 1162.07




Test Example (idx: 195)
True: 3
Predicted: 1 (Misclassified)




Neighbor 1
Label: 5
Dist: 1368.12




Neighbor 2
Label: 5
Dist: 1484.98




Neighbor 3
Label: 1
Dist: 1488.17



Neighbor 4
Label: 1
Dist: 1491.23



Neighbor 5
Label: 1
Dist: 1503.51






Choosing k

k vs. # of Misclassifications


k	1	3	5	7	9	11	13	15
FP + FN	309	295	318	306	341	332	347	367

k = 3, as it resulted in the least number of misclassifications.


Test Example (idx: 33)
True: 4
Predicted: 0 (Misclassified)




Neighbor 1
Label: 4
Dist: 1384.31




Neighbor 2
Label: 0
Dist: 1780.64




Neighbor 3
Label: 4
Dist: 1841.15



Neighbor 4
Label: 0
Dist: 1897.24



Neighbor 5
Label: 0
Dist: 1904.74



k = 5

MNIST is pre-split into 60,000 Lines of Training and 10,000 Testing.




Choosing k

k vs. # of Misclassifications


k	1	3	5	7	9	11	13	15
FP + FN	309	295	318	306	341	332	347	367

k = 3, as it resulted in the least number of misclassifications.


Test Example (idx: 33)
True: 4
Predicted: 0 (Misclassified)




Neighbor 1
Label: 4
Dist: 1384.31



Neighbor 2
Label: 0
Dist: 1780.64



Neighbor 3
Label: 4
Dist: 1841.15




k = 3

MNIST is pre-split into 60,000 Lines of Training and 10,000 Testing.



Validation Metrics



Macro-averaging metrics is a method of getting validation metrics for **non-binary classification** by taking the **mean** of the metrics for every **individual class**.

Macro-Average Metrics:

- Precision: **0.9709**
- F1: **0.9701**
- Recall: **0.9704**
- Accuracy: **0.9941**



05

Results

Evaluating Results



Results

- KNN can transform images into text at **99.41% accuracy**.
- Can be scaled through increasing size of training data with more variations.
 - Can also be expanded to interpret letters or special characters.

a	B	C	d	e	F	G	H	I	J	K	L	M	N	O	P	q	R	S	T	U	V	W	X	Y	Z	0	0
A	b	C	D	e	f	G	H	I	J	K	L	M	N	O	P	q	R	S	T	U	V	W	X	Y	Z	1	1
a	b	C	D	e	f	G	H	I	J	K	L	M	N	O	P	q	R	S	T	U	V	W	X	Y	Z	2	2
a	b	C	D	e	f	G	H	I	J	K	L	M	N	O	P	q	R	S	T	U	V	W	X	Y	Z	3	3
A	B	C	d	e	f	G	H	I	J	K	L	M	N	O	P	q	R	S	T	U	V	W	X	Y	Z	4	4
a	B	C	d	e	f	G	H	I	J	K	L	M	N	O	P	q	R	S	T	U	V	W	X	Y	Z	5	5
A	B	C	d	e	f	G	H	I	J	K	L	M	N	O	P	q	R	S	T	U	V	W	X	Y	Z	6	6
a	b	C	d	e	f	G	H	I	J	K	L	M	N	O	P	q	R	S	T	U	V	W	X	Y	Z	7	7
a	b	C	d	e	f	G	H	I	J	K	L	M	N	O	P	q	R	S	T	U	V	W	X	Y	Z	8	8
A	B	C	d	e	f	G	H	I	J	K	L	M	N	O	P	q	R	S	T	U	V	W	X	Y	Z	9	9



What I Learned?

KNN proves effective for image-to-text conversion, offering a foundation for developing more sophisticated OCR systems.

Each k-value test took 20-30 minutes to process 10,000 test images, varying based on whether the program ran in the foreground or background of my system.

Real World Applications:

- Document Processing
- Data Entry Automation
- Accessibility Tools

Limitations of KNN for OCR:

- Computationally Expensive: $O(nd)$
- Scaling: Slower with Larger Datasets
- Memory Requirements: Stores Training Data

Though KNN is effective for OCR it is nowhere near as practical as the more widely used alternative, **Convolutional Neural Networks**.



Thank You

Any Questions?



CREDITS: This presentation template was created by
Slidesgo, and includes icons by **Flaticon**, and infographics
& images by **Freepik**